

Predicting activity(exercise) quality from activity monitors using machine learning algorithm

jshiju

Saturday, August 22, 2015

```
## Run time: 2015-08-22 18:43:29
## R version: R version 3.2.0 (2015-04-16)
## OS: Windows 7 x64
```

I. SYNOPSIS:

The purpose of this project is to develop and build a model using machine learning techniques, based on the WLE(Weight Lifting Exercise) Dataset, to predict the manner in which an health participant performed an exercise on 20 different test cases with 'classe' as the resspense variable.

II. DATASET & DESCRIPTION:

The WLE Dataset is available at <http://groupware.les.inf.puc-rio.br/har> and was collected from sensors(accelerometers) on the belt, forearm, arm, and dumbbell of Six health participants) who were asked to perform one set of 10 repetitions of the Unilateral Dumbbell Biceps Curl in five different fashions: exactly according to the specification (Class A), throwing the elbows to the front (Class B), lifting the dumbbell only halfway (Class C), lowering the dumbbell only halfway (Class D) and throwing the hips to the front (Class E). Class A corresponds to the specified execution of the exercise, while the other 4 classes correspond to common mistakes.

Training data : <https://d396qusza40orc.cloudfront.net/predmachlearn/pml-training.csv> Test data : <https://d396qusza40orc.cloudfront.net/testing.csv>

III. DATA PREPARATION

Read training and test datasets from the source

```
# read train data set
require(data.table)
setInternet2(TRUE)
url <- "https://d396qusza40orc.cloudfront.net/predmachlearn/pml-training.csv"
pmlTrainData <- fread(url)

##
Read 51.0% of 19622 rows
Read 19622 rows and 160 (of 160) columns from 0.011 GB file in 00:00:04

# read test data set
url <- "https://d396qusza40orc.cloudfront.net/predmachlearn/pml-testing.csv"
pmlTestData <- fread(url)
```

IV. EDA & PREDICTOR IDENTIFICATION

Perform an EDA(Exploratory Data Analysis) on the data set.

```
#summary(pmlTrainData)
#describe(pmlTrainData)
#sapply(pmlTrainData, class)
#str(pmlTrainData)
```

A quick analysis on the test data set shows that we cannot take into all the variables for the prediction and need to identify those predictor variables which are relevant. We are interested in those variables produced by sensors with a Non-NA values.

Also subset the primary dataset to include only the predictor candidates and the outcome/response variable - 'classe'.

```
# identify predictors
isAnyMissing <- sapply(pmlTestData, function(x) any(is.na(x) | x == ""))
isPredictor <- !isAnyMissing & grepl("belt|^(fore)|arm|dumbbell|forearm", names(isAnyMissing))
predCandidates <- names(isAnyMissing)[isPredictor]

# subset primary dataset for predictor & outcome variables
varToInclude <- c("classe", predCandidates)
pmlTrainData <- pmlTrainData[, varToInclude, with=FALSE]
```

Perform the required Data Cleansing operations and split the dataset into training and probing dataset in the ratio 60:40. And a final look at the dataset attributes

```
# classe as factor
pmlTrainData <- pmlTrainData[, classe := factor(pmlTrainData[, classe])]
dim(pmlTrainData)

## [1] 19622    53

names(pmlTrainData)

## [1] "classe"          "roll_belt"       "pitch_belt"
## [4] "yaw_belt"        "total_accel_belt" "gyros_belt_x"
## [7] "gyros_belt_y"    "gyros_belt_z"    "accel_belt_x"
## [10] "accel_belt_y"    "accel_belt_z"    "magnet_belt_x"
## [13] "magnet_belt_y"   "magnet_belt_z"   "roll_arm"
## [16] "pitch_arm"       "yaw_arm"         "total_accel_arm"
## [19] "gyros_arm_x"     "gyros_arm_y"     "gyros_arm_z"
## [22] "accel_arm_x"     "accel_arm_y"     "accel_arm_z"
## [25] "magnet_arm_x"    "magnet_arm_y"    "magnet_arm_z"
## [28] "roll_dumbbell"   "pitch_dumbbell"  "yaw_dumbbell"
## [31] "total_accel_dumbbell" "gyros_dumbbell_x" "gyros_dumbbell_y"
## [34] "gyros_dumbbell_z" "accel_dumbbell_x" "accel_dumbbell_y"
## [37] "accel_dumbbell_z" "magnet_dumbbell_x" "magnet_dumbbell_y"
## [40] "magnet_dumbbell_z" "roll_forearm"    "pitch_forearm"
## [43] "yaw_forearm"     "total_accel_forearm" "gyros_forearm_x"
## [46] "gyros_forearm_y" "gyros_forearm_z"  "accel_forearm_x"
## [49] "accel_forearm_y" "accel_forearm_z"  "magnet_forearm_x"
## [52] "magnet_forearm_y" "magnet_forearm_z"
```

```
pmlTrainData[, .N, classe]
```

```
##   classe    N
## 1:    A 5580
## 2:    B 3797
## 3:    C 3422
## 4:    D 3216
## 5:    E 3607
```

```
# split dataset [60% - training; 40% - probing]
require(caret)
seed <- as.numeric(as.Date("2015-08-21"))
set.seed(seed)
inTrain <- createDataPartition(pmlTrainData$classe, p=0.6)
trainData <- pmlTrainData[inTrain[[1]]]
probeData <- pmlTrainData[-inTrain[[1]]]
```

The next step would be to estimate pre-processing transformation (centering, scaling etc) from the training data and applied to probe data set with the same variables. Also diagnose predictors for near zero variance.

```
# preprocess the prediction variables by centering and scaling.
origData <- trainData[, predCandidates, with=FALSE]
preProcessor <- preProcess(origData)
tranformData <- predict(preProcessor, origData)
DTrainCS <- data.table(data.frame(classe = trainData[, classe], tranformData))

# apply the centering and scaling to the probing dataset.
origData <- probeData[, predCandidates, with=FALSE]
tranformData <- predict(preProcessor, origData)
DProbeCS <- data.table(data.frame(classe = probeData[, classe], tranformData))

# check for near zero variance.
nzv <- nearZeroVar(DTrainCS, saveMetrics=TRUE)
if (any(nzv$nzv)) nzv else message("No variables with near zero variance")

## No variables with near zero variance
```

Examine groups of prediction variables and its replationship with response variable using plotting.

```
require(reshape2)
require(ggplot2)

histGroup <- function (data, regex) {
  col <- grep(regex, names(data))
  col <- c(col, which(names(data) == "classe"))
  n <- nrow(data)
  DMelted <- melt(data[, col, with=FALSE][, rownum := seq(1, n)], id.vars=c("rownum", "classe"))

  ggplot(DMelted, aes(x=classe, y=value)) +
    geom_violin(aes(color=classe, fill=classe), alpha=1/2) +
    facet_wrap(~ variable, scale="free_y") +
    scale_color_brewer(palette="Spectral") +
    scale_fill_brewer(palette="Spectral") +
    labs(x="", y="") +
    theme(legend.position="none")
}

histGroup(DTrainCS, "belt")

histGroup(DTrainCS, "[^(fore)]arm")

histGroup(DTrainCS, "dumbbell")

histGroup(DTrainCS, "forearm")
```

V. FITTING A MODEL USING RANDOM FOREST

```
# set up the parallel clusters.
require(parallel)
require(doParallel)
cl <- makeCluster(detectCores() - 1)
registerDoParallel(cl)
```

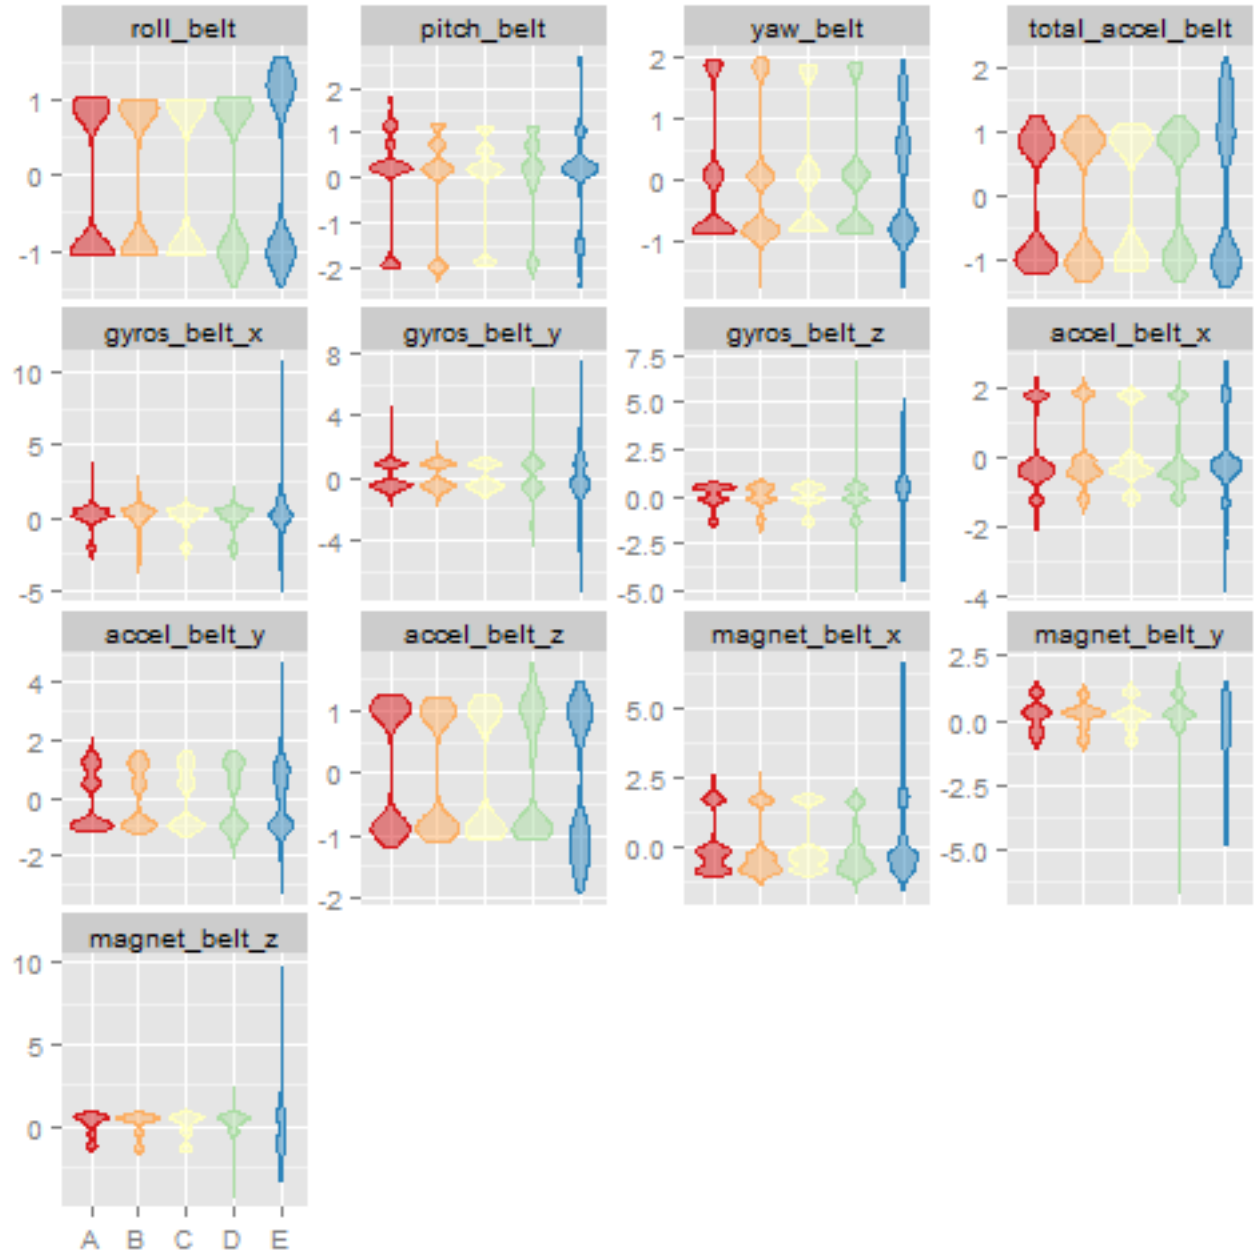


Figure 1: plot of chunk unnamed-chunk-8

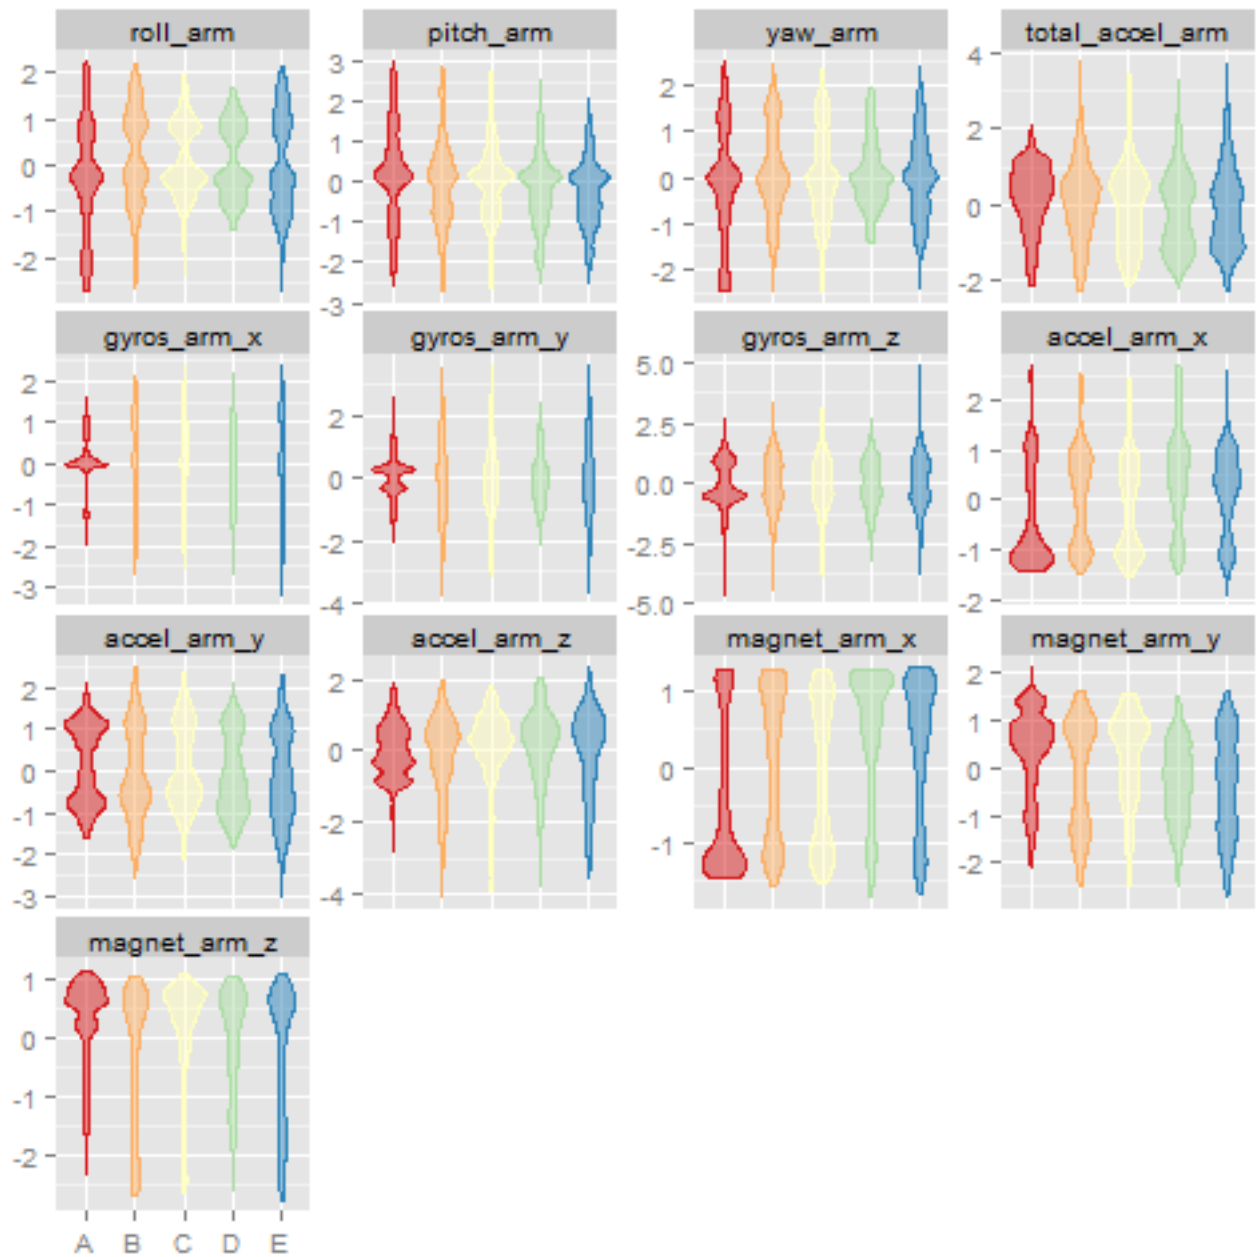


Figure 2: plot of chunk unnamed-chunk-8

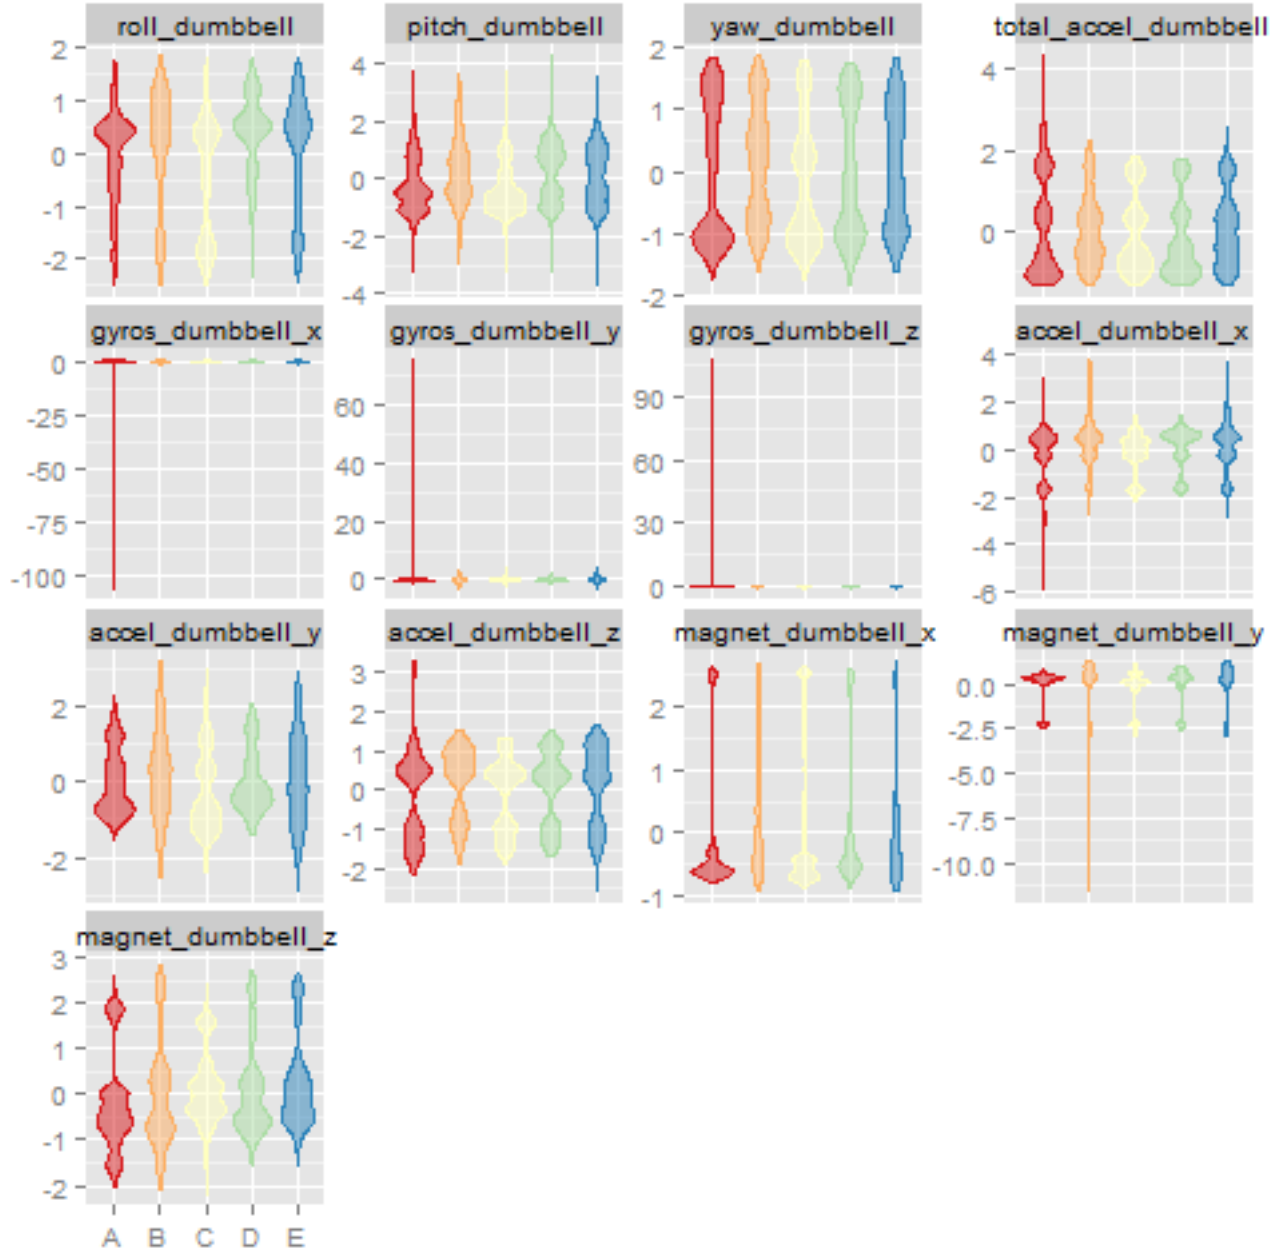


Figure 3: plot of chunk unnamed-chunk-8

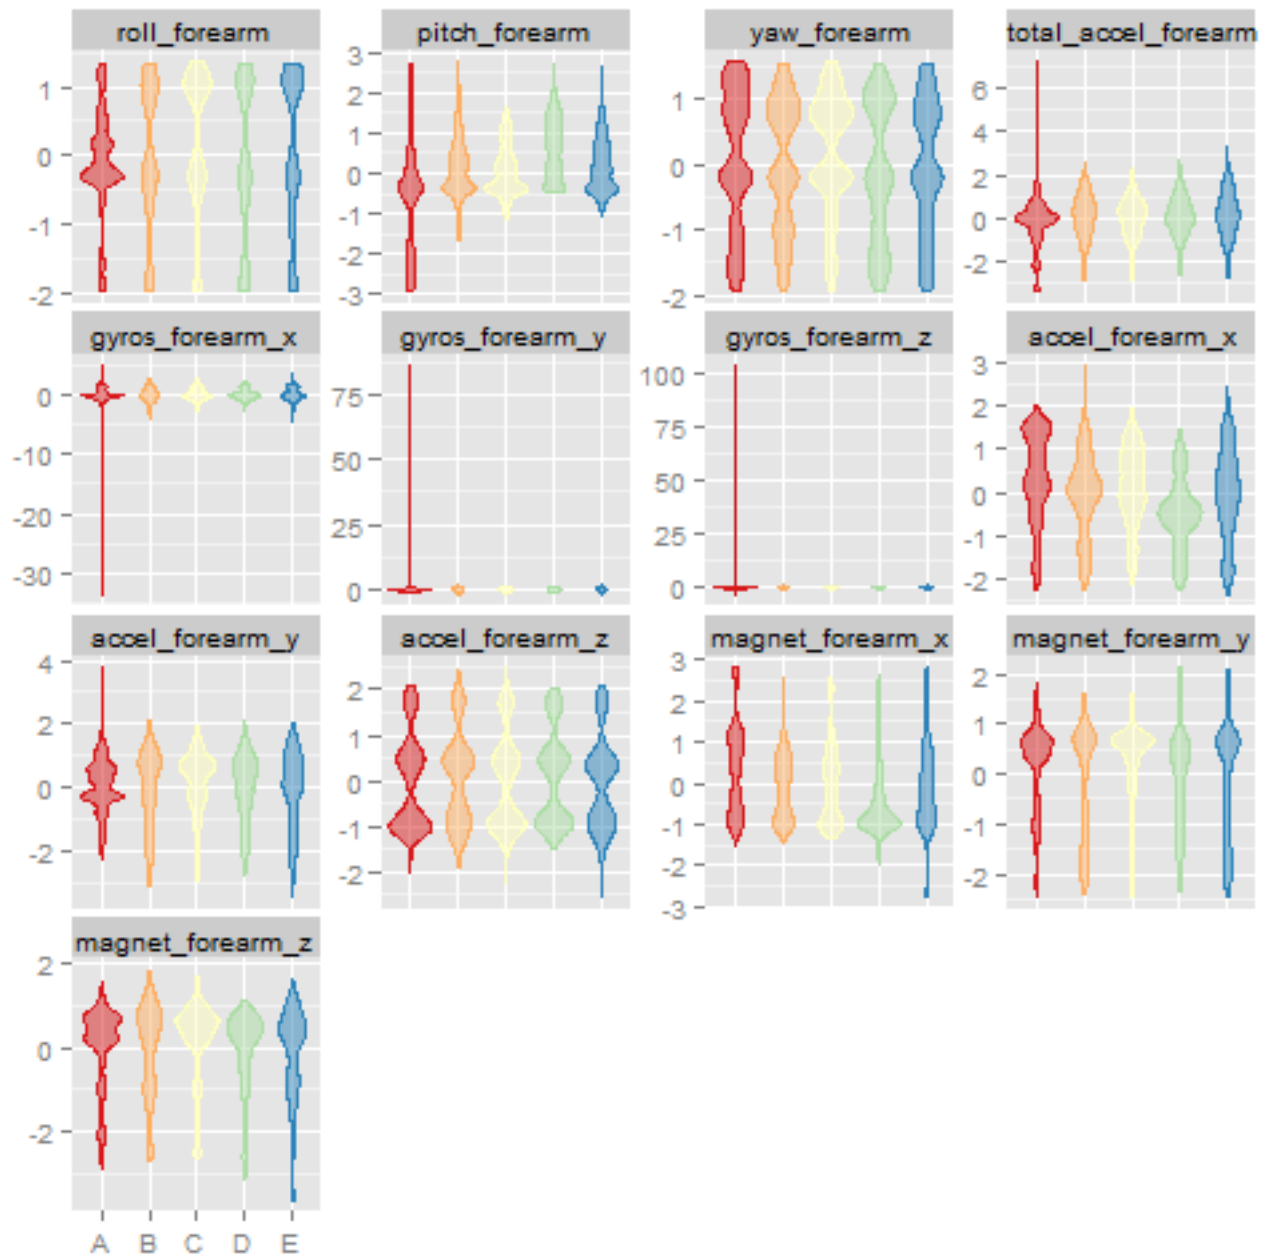


Figure 4: plot of chunk unnamed-chunk-8

```
# set the control parameters.
ctrl <- trainControl(classProbs=TRUE,
                      savePredictions=TRUE,
                      allowParallel=TRUE)

#fit model (random forrest) over the tuning parameters.
trainingModel <- train(classe ~ ., data=DTrainCS, method="rf")
#system.time(trainingModel <- train(classe ~ ., data=DTrainCS, method="rf"))

# stop the clusters.
stopCluster(cl)
```

VI. OUT-OF SAMPLE ERROR & ERROR ESTIMATE

To evaluate the model we will use the confusionmatrix method and we will focus on accuracy, sensitivity & specificity metrics. As seen from the result of the confusionmatrix below, the model is good and efficient because it has an accuracy of 0.997 and very good sensitivity & specificity values on the testing dataset. (the lowest value is 0.992 for the sensitivity of the class C)

The estimated error rate is less than 1%.

```
# evaluate the model on the training dataset
trainingModel

## Random Forest
##
## 11776 samples
## 52 predictor
## 5 classes: 'A', 'B', 'C', 'D', 'E'
##
## No pre-processing
## Resampling: Bootstrapped (25 reps)
## Summary of sample sizes: 11776, 11776, 11776, 11776, 11776, ...
## Resampling results across tuning parameters:
##
## mtry Accuracy Kappa Accuracy SD Kappa SD
## 2 0.9860429 0.9823473 0.001776708 0.002240671
## 27 0.9860108 0.9823071 0.001742289 0.002198113
## 52 0.9798819 0.9745544 0.003910534 0.004951608
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was mtry = 2.
```

```
hat <- predict(trainingModel, DTrainCS)
confusionMatrix(hat, trainData[, classe])
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  A    B    C    D    E
## A 3348      0     0     0     0
## B  0 2279      0     0     0
## C  0  0 2054      0     0
## D  0  0  0 1930      0
## E  0  0  0  0 2165
##
## Overall Statistics
##
##           Accuracy : 1
##           95% CI : (0.9997, 1)
## No Information Rate : 0.2843
## P-Value [Acc > NIR] : < 2.2e-16
```



```
##
##           Kappa : 1
##  McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity      1.0000  1.0000  1.0000  1.0000  1.0000
## Specificity      1.0000  1.0000  1.0000  1.0000  1.0000
## Pos Pred Value   1.0000  1.0000  1.0000  1.0000  1.0000
## Neg Pred Value   1.0000  1.0000  1.0000  1.0000  1.0000
## Prevalence       0.2843  0.1935  0.1744  0.1639  0.1838
## Detection Rate   0.2843  0.1935  0.1744  0.1639  0.1838
## Detection Prevalence 0.2843  0.1935  0.1744  0.1639  0.1838
## Balanced Accuracy 1.0000  1.0000  1.0000  1.0000  1.0000
```

```
# final model
varImp(trainingModel)
```

```
## rf variable importance
##
##   only 20 most important variables shown (out of 52)
##
##           Overall
## roll_belt      100.00
## yaw_belt       78.32
## magnet_dumbbell_z 65.79
## magnet_dumbbell_y 58.90
## pitch_belt     58.65
## pitch_forearm  58.00
## roll_forearm   48.64
## magnet_dumbbell_x 47.05
## magnet_belt_z  42.79
## accel_belt_z   41.64
## accel_dumbbell_y 41.63
## magnet_belt_y  41.44
## roll_dumbbell  40.76
## accel_dumbbell_z 37.55
## roll_arm       34.26
## gyros_belt_z   31.72
## accel_forearm_x 31.20
## yaw_dumbbell   28.53
## total_accel_dumbbell 27.42
## accel_dumbbell_x 26.64
```

```
trainingModel$finalModel
```

```
##
## Call:
##  randomForest(x = x, y = y, mtry = param$mtry)
##           Type of random forest: classification
##           Number of trees: 500
## No. of variables tried at each split: 2
##
##           OOB estimate of  error rate: 0.94%
## Confusion matrix:
##           A      B      C      D      E  class.error
## A 3346      0      1      0      1 0.0005973716
## B   23 2247      9      0      0 0.0140412462
## C    0   22 2029      3      0 0.0121713729
```

```
## D    0    0   39 1887    4 0.0222797927
## E    0    0    2    7 2156 0.0041570439
```

```
# save training model
save(trainingModel, file="trainingModel.RData")
```

```
# Get predictions and evaluate.
```

```
DTestCS <- predict(preProcessor, pmlTestData[, predCandidates, with=FALSE])
```

```
hat <- predict(trainingModel, DTestCS)
```

```
pmlTestData <- cbind(hat , pmlTestData)
```

```
subset(pmlTestData, select=names(pmlTestData)[grep("belt|[^(fore)]arm|dumbbell|forearm", names(pmlTestData), inv
```

```
##      hat V1 user_name raw_timestamp_part_1 raw_timestamp_part_2
## 1:  B  1   pedro      1323095002      868349
## 2:  A  2   jeremy      1322673067      778725
## 3:  B  3   jeremy      1322673075      342967
## 4:  A  4   adelmo      1322832789      560311
## 5:  A  5   eurico      1322489635      814776
## 6:  E  6   jeremy      1322673149      510661
## 7:  D  7   jeremy      1322673128      766645
## 8:  B  8   jeremy      1322673076       54671
## 9:  A  9  carlitos      1323084240      916313
## 10: A 10  charles      1322837822      384285
## 11: B 11  carlitos      1323084277       36553
## 12: C 12   jeremy      1322673101      442731
## 13: B 13   eurico      1322489661      298656
## 14: A 14   jeremy      1322673043      178652
## 15: E 15   jeremy      1322673156      550750
## 16: E 16   eurico      1322489713      706637
## 17: A 17   pedro      1323094971      920315
## 18: B 18  carlitos      1323084285      176314
## 19: B 19   pedro      1323094999      828379
## 20: B 20   eurico      1322489658      106658
```

```
##      cvtd_timestamp new_window num_window problem_id
## 1: 05/12/2011 14:23      no         74          1
## 2: 30/11/2011 17:11      no        431          2
## 3: 30/11/2011 17:11      no        439          3
## 4: 02/12/2011 13:33      no        194          4
## 5: 28/11/2011 14:13      no        235          5
## 6: 30/11/2011 17:12      no        504          6
## 7: 30/11/2011 17:12      no        485          7
## 8: 30/11/2011 17:11      no        440          8
## 9: 05/12/2011 11:24      no        323          9
## 10: 02/12/2011 14:57      no        664         10
## 11: 05/12/2011 11:24      no        859         11
## 12: 30/11/2011 17:11      no        461         12
## 13: 28/11/2011 14:14      no        257         13
## 14: 30/11/2011 17:10      no        408         14
## 15: 30/11/2011 17:12      no        779         15
## 16: 28/11/2011 14:15      no        302         16
## 17: 05/12/2011 14:22      no         48         17
## 18: 05/12/2011 11:24      no        361         18
## 19: 05/12/2011 14:23      no         72         19
## 20: 28/11/2011 14:14      no        255         20
```

VII. PREDICTION ON TEST SET

And the final prediction on the 20 test cases are:

```
1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 B A B A A E D B A A B C B A E E A B B B
```